


# Chapter 9: Virtual-Memory Management

---





# Virtual Memory

---

## Goal

Provide memory space for all processes that can be much larger than physical memory.

## How

Only some part of program code and data need to be in physical memory at a time.

- ▶ principle of locality

The rest of program code and data can be stored in secondary storage (hard disk)

- ▶ “swap partition” or “swap file”
- ▶ extend memory hierarchy

Implement “Demand Paging” mechanism in OS

- ▶ similar to “cache miss”

## Other benefits

Address spaces can be shared by several processes

- ▶ shared program code, shared libraries, shared memory (IPC)


fast process creation (allocate memory only when needed)

memory-mapped file, memory-mapped I/O





# Virtual Memory





# Demand Paging

---

Bring a page into memory only when it is needed

- Less I/O needed

- Less memory needed

- Faster response

- More processes

CPU instruction references to anywhere in a page  $\Rightarrow$  the page needs to be in physical memory

- invalid reference  $\Rightarrow$  abort


- not-in-memory  $\Rightarrow$  bring to memory

- no-free-frame  $\Rightarrow$  swap





# Page table that supports demand paging





# Page Fault

---

Page fault is a mechanism to implement Demand Paging

Similar to cache miss

A reference to an invalid page (e.g. first reference to that page) will cause an interrupt to operating system:

page fault interrupt (generated by MMU)

page fault handler (code in OS kernel)

1. Operating system looks at another table to decide:

Invalid reference → abort

Just not in memory → continue

2. Get empty frame

3. Swap page into frame

4. Update tables


5. Set validation bit = **v**

6. Restart the instruction that caused the page fault





# Steps in Handling a Page Fault






# Copy-on-Write


**Copy-on-Write** (COW) allows both parent and child processes to initially *share* the same pages in memory

If either process modifies a shared page, only then is the page copied

COW allows more efficient process creation as only modified pages are copied



Before Process 1 Modifies Page C




After Process 1 Modifies Page C







# Need For Page Replacement





# Basic Page Replacement


---

1. Find the location of the desired page on disk.
2. Find a free frame
  - If there is a free frame, use it.
  - If there is no free frame, use a page replacement algorithm to select a **victim frame**.
  - Write victim frame to disk if dirty.
3. Bring the desired page into the (newly) free frame; update the page and frame tables.
4. Continue the process by restarting the instruction that caused the trap.





# Page Replacement





# First-In-First-Out (FIFO) Algorithm

Use a FIFO queue. When a page is brought into memory, append it into the queue.


When a page must be replaced, choose the oldest page in memory (the first one in the queue).

Example: Reference string: **7,0,1,2,0,3,0,4,2,3,0,3,2,1,2,0,1,7,0,1**

3 frames (3 pages can be in memory at a time per process)

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1




page frames

15 page faults

Adding more frames can cause more page faults!

**Belady's Anomaly**

consider reference string: 1,2,3,4,1,2,5,1,2,3,4,5





# Optimal Algorithm

Replace page that will not be used for longest period of time

How do you know this?

Can't read the future

Used for measuring how well your algorithm performs

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1



page frames

9 page faults

Reference to page 2 causes a page replacement  
7 will be used again in 14 references  
0 will be used again in 1 reference  
1 will be used again in 10 reference  
So, replace 7





# Least Recently Used (LRU) Algorithm


Use past knowledge rather than future

Replace page that has not been used in the most amount of time

Associate time of last use with each page

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1



page frames

Reference to page 2 causes a page replacement  
7 has not been used for 3 references  
0 has not been used for 2 references  
1 has not been used for 1 references  
So, replace 7

12 faults – better than FIFO but worse than Optimal Algorithm

Generally good algorithm and frequently used

But how to implement?





# LRU Algorithm (Cont.)

## Counter implementation

Every page entry has a counter; every time page is referenced through this entry, copy the clock into the counter

When a page needs to be changed, look at the counters to find smallest value

- ▶ Search through table needed

## Stack implementation

Keep a stack of page numbers in a double link form:

Page referenced:

- ▶ move it to the top
- ▶ requires 6 pointers to be changed

But each update more expensive

No search for replacement

LRU and OPT are cases of **stack algorithms** that don't have Belady's Anomaly

In a stack algorithm, the set of pages in memory for  $n$  frames is always a subset of the set of pages that would be in memory with  $n+1$  frames.





# LRU Approximation Algorithms

LRU needs special hardware and still slow

## Reference bit

With each page associate a bit, initially = 0

When page is referenced bit set to 1

Replace any with reference bit = 0 (if one exists)

- ▶ We do not know the order, however

## Second-chance algorithm

Generally FIFO, plus hardware-provided reference bit

Clock replacement

If page to be replaced has


- ▶ Reference bit = 0 -> replace it
- ▶ reference bit = 1 then:
  - set reference bit 0, leave page in memory
  - replace next page, subject to same rules







## Second-Chance (clock) Page-Replacement Algorithm



Use a circular queue to keep pages in FIFO manner.

- Oldest page in memory is at the head of queue and will be replaced first.
- But if it has been used, it will be given a second chance.
- Each page has a reference bit. Initialized to zero and set to one when accessed. Also reset to zero periodically.

If page to be replaced has

- Reference bit = 0 → replace it
- reference bit = 1 then:
  - ▶ set reference bit 0, leave page in memory
  - ▶ replace next page, subject to same rules

Note: Second-Chance algorithm can behave like FIFO if all pages are frequently accessed so they all get the second chance and finally the first page in the queue is the victim. Normally, this is unlikely.





# Shared Pages

---


- Virtual memory allows pages to be shared among processes
- Logical memory pages of multiple processes are mapped to the same physical memory frames.
- This technique enables many capabilities:
  - Shared libraries
  - Memory-Mapped files
  - Shared-memory communication





# Shared Library Using Virtual Memory

- Code libraries can be shared by several processes.
- Each process considers the library as part of its virtual address space.
- The actual pages where the libraries reside are shared.





# Memory-Mapped Files

---

Memory-mapped file I/O allows file I/O to be treated as routine memory access by **mapping** a disk block to a page in memory

A file is initially read using demand paging

A page-sized portion of the file is read from the file system into a physical page

Subsequent reads/writes to/from the file are treated as ordinary memory accesses

Simplifies and speeds file access by driving file I/O through memory rather than `read()` and `write()` system calls

Also allows several processes to map the same file allowing the pages in memory to be shared


But when does written data make it to disk?

Periodically and / or at file `close()` time





# Memory Mapped Files





# Memory-Mapped Shared Memory

- Processes can communicate through shared memory.
- A process can create a region of memory that it share with another process.
- Processes sharing this region consider it part of their virtual address space, yet the actual physical pages of memory are shared.

